

# Pythran: Static Compilation of Parallel Scientific Kernels

a.k.a. Python/Numpy compiler for the  
mass

Proudly made in *Namek* by serge-sans-paille & pbrunet

# /us

## Serge « sans paille » Guelton

```
$ whoami  
sguelton
```


- R&D engineer at QuarksLab on compilation for security
- Associate researcher at Télécom Bretagne

## Pierrick Brunet

```
$ whoami  
pbrunet
```

- R&D engineer at INRIAAlpes/MOAIS on parallelism

# Pythran in a snake shell



- A Numpy-centric Python-to-C++ translator
- A Python code *optimizer*
- A Pythonic C++ library

# Core concepts

- Focus on high-level constructs
- Generate ~~clean~~ high level code
- Optimize Python code before generated code
- Vectorization and Parallelism
- Test, test, test
- Bench, bench, bench

# Ask StackOverflow

*when you're looking for **test cases***

[http://stackoverflow.com/\[...\]numba-or-cython-acceleration-in-reaction-diffusion-algorithm](http://stackoverflow.com/[...]numba-or-cython-acceleration-in-reaction-diffusion-algorithm)

```
import numpy as np
def GrayScott(counts, Du, Dv, F, k):
    n = 300
    U = np.zeros((n+2,n+2), dtype=np.float32)
    V = np.zeros((n+2,n+2), dtype=np.float32)
    u, v = U[1:-1,1:-1], V[1:-1,1:-1]

    r = 20
    u[:] = 1.0
    U[n/2-r:n/2+r,n/2-r:n/2+r] = 0.50
    V[n/2-r:n/2+r,n/2-r:n/2+r] = 0.25
    u += 0.15*np.random.random((n,n))
    v += 0.15*np.random.random((n,n))

    for i in range(counts):
        Lu = (
            U[0:-2,1:-1] +
            U[1:-1,0:-2] - 4*U[1:-1,1:-1] + U[1:-1,2:] +
            U[2:,1:-1])
```

# Thread Summary

**OP**

My code is slow with Cython and Numba

**Best Answer**

You need to make all loops explicit

# Cython Version

```
cimport cython
import numpy as np
cimport numpy as np

cpdef cythonGrayScott(int counts, double Du, double Dv, double F, double
    cdef int n = 300
    cdef np.ndarray U = np.zeros((n+2,n+2), dtype=np.float_)
    cdef np.ndarray V = np.zeros((n+2,n+2), dtype=np.float_)
    cdef np.ndarray u = U[1:-1,1:-1]
    cdef np.ndarray v = V[1:-1,1:-1]

    cdef int r = 20
    u[:] = 1.0
    U[n/2-r:n/2+r,n/2-r:n/2+r] = 0.50
    V[n/2-r:n/2+r,n/2-r:n/2+r] = 0.25
    u += 0.15*np.random.random((n,n))
    v += 0.15*np.random.random((n,n))
```

# Pythran version

Add this line to the original kernel:

```
#pythran export GrayScott(int, float, float, float, float)
```

## Timings

```
$ python -m timeit -s 'from grayscott import GrayScott' 'GrayScott(40, 0, 0, 0, 0)'
10 loops, best of 3: 52.9 msec per loop
$ cython grayscott.pyx
$ gcc grayscott.c `python-config --cflags --libs` -shared -fPIC -o grayscott.so
$ python -m timeit -s 'from grayscott import GrayScott' 'GrayScott(40, 0, 0, 0, 0)'
10 loops, best of 3: 36.4 msec per loop
$ pythran grayscott.py -O3 -march=native
$ python -m timeit -s 'from grayscott import GrayScott' 'GrayScott(40, 0, 0, 0, 0)'
10 loops, best of 3: 20.3 msec per loop
```



# Lessons learnt

- Explicit is not always better than implicit
- Many "optimization hints" can be deduced by the compiler
- High level constructs carry valuable informations

*I am **not** saying Cython is bad. Cython does a **great** job. It is just **pragmatic** where Pythran is **idealist***

# Compilation Challenges

```
u = U[1:-1,1:-1]  
U[n/2-r:n/2+r,n/2-r:n/2+r] = 0.50  
u += 0.15*np.random.random((n,n))
```

- Array views
- Value broadcasting
- Temporary arrays creation
- Extended slices composition
- Numpy API calls

# Optimization Opportunities

```
Lu = (U[0:-2,1:-1] + U[1:-1,0:-2]  
      - 4*U[1:-1,1:-1] + U[1:-1,2:] + U[2: ,1:-1])
```

- Many useless temporaries
- Lu could be forward-substituted
- SIMD instruction generation opportunities
- Parallel loop opportunities

# Pythran Usage

```
$ pythran --help
usage: pythran [-h] [-o OUTPUT_FILE] [-E] [-e] [-f flag] [-v] [-p pass]
              [-m machine] [-I include_dir] [-L ldflags]
              [-D macro_definition] [-O level] [-g]
              input_file

pythran: a python to C++ compiler

positional arguments:
  input_file            the pythran module to compile, either a .py or a
                        file

optional arguments:
  -h, --help            show this help message and exit
  -o OUTPUT_FILE        path to generated file
  -E                    only run the translator, do not compile
  -e                    similar to -E, but does not generate python glue
  -f flag               any compiler switch relevant to the underlying C++
```

# Sample Usage

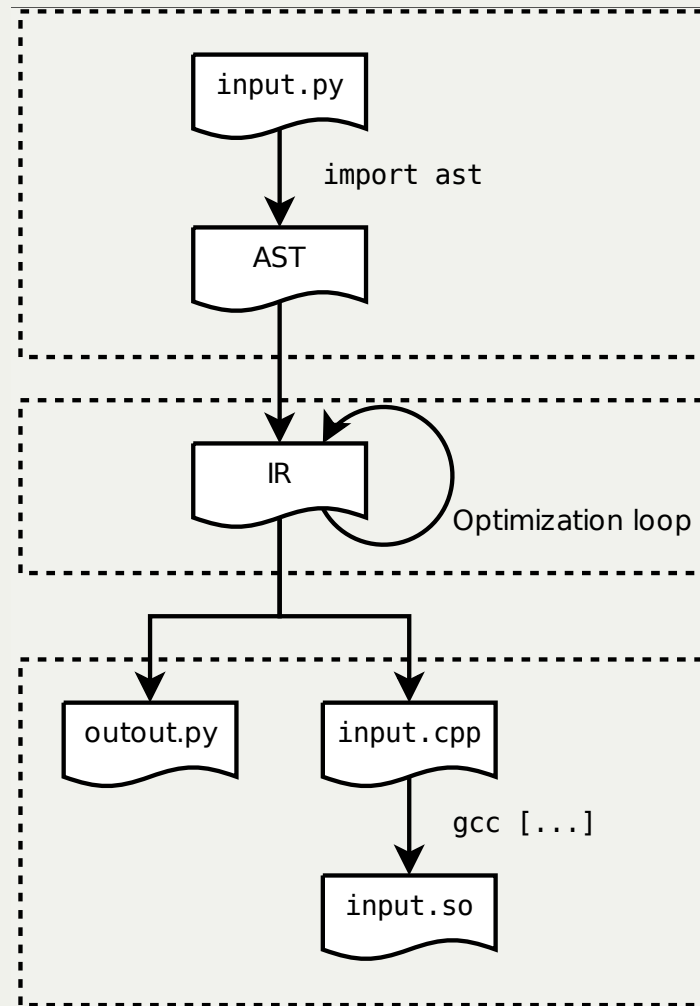
```
$ pythran input.py # generates input.so  
$ pythran input.py -E # generates input.cpp  
$ pythran input.py -O3 -fopenmp # parallel!  
$ pythran input.py -march=native -Ofast # Esod Mumixam !
```

# Type Annotations

Only for *exported* functions

```
#pythran export foo0()  
#pythran export foo1(int)  
  
#pythran export foo2(float32[][])  
#pythran export foo2(float64[][])  
#pythran export foo2(int8[] [] [])  
  
#pythran export foo3((int, float), int list, str:str dict)
```

# Pythran Compilation Flow



# Front End

- 100% based on the `ast` module
- Supports
  - Several standard module (incl. partial Numpy)
  - Polymorphic functions
  - ndarray, list, tuple, dict, str, int, long, float
  - Named parameters, default arguments
  - Generators...
- Does not Support
  - Non-implicitly typed code
  - Global variable
  - Most Python modules (no CPython mode!)
  - User-defined classes...



# Middle End

Iteratively applies high level, Python-aware optimizations:

- Interprocedural Constant Folding
- For-Based-Loop Unrolling
- Forward Substitution
- Instruction Selection
- Deforestation
- Scalar Renaming
- dead Code Elimination

*Fun Facts: can evaluate pure functions at compile time ☺*

# Back Ends

## Python Back End

Useful for debugging!

## C++11 Back End

- C++11 implementation of `__builtin__` numpy `itertools...`
- Lazy evaluation through Expression Templates
- Relies on OpenMP, `nt2` and `boost::simd` for the parallelization / vectorization

# C++11: Typing

*the W.T.F. slide*

Pythran translates Python implicitly statically typed **polymorphic** code into C++ **meta-programs** that are instantiated for the user-given types, and specialize them for the target architecture

# C++11: Parallelism

## Implicit

Array operations and several numpy functions are written using  
OpenMP and Boost.simd

## Explicit

OpenMP 3 support, ported to Python

```
#omp parallel for reduction(+:r)
for i, v in enumerate(l):
    r += i * v
```

# Benchmarks

<https://github.com/serge-sans-paille/numpy-benchmarks>


A collection of high-level benchmarks

- Code gathered from StackOverflow + other compiler code base
- Mostly high-level code
- Generate results for CPython, PyPy, Numba, Parakeet, Hope and Pythran

*Most kernels are too high level for Numba and Hope...*

# Benchmarks

no parallelism, no vectorisation (, no fat)



# (Num)Focus: growcut

From the Numba codebase!

```
#pythran export growcut(float[][][], float[][][], float[][][], int)
import math
import numpy as np


def window_floor(idx, radius):
    if radius > idx:
        return 0
    else:
        return idx - radius

def window_ceil(idx, ceil, radius):
    if idx + radius > ceil:
        return ceil
    else:
        return idx + radius

def growcut(image, state, state_next, window_radius):
    changes = 0
```



# (Num)Focus: growcut



# Academic Results

- **Pythran: Enabling Static Optimization of Scientific Python Programs**, S. Guelton, P. Brunet et al. in *CSD*, 2015
- **Exploring the Vectorization of Python Constructs Using Pythran and Boost SIMD**, S. Guelton, J. Falcou and P. Brunet, in *WPMVP*, 2014
- **Compiling Python modules to native parallel modules using Pythran and OpenMP Annotations**, S. Guelton, P. Brunet and M. Amini, in *PyHPC*, 2013
- **Pythran: Enabling Static Optimization of Scientific Python Programs**, S. Guelton, P. Brunet et al. in *SciPy*, 2013

# Powered by Strong Engineering

Preprerequisite for reproducible science

- 2773 test cases, incl. unit testing, doctest, **Continuous integration** (thx Travis!)
- Peer-reviewed code
- Python2.7 and C++11
- Linux, OSX (almost okay), Windows (on going)
- User and Developer doc: <http://pythonhosted.org/pythran/>
- Hosted on <https://github.com/serge-sans-paille/pythran>
- Releases on PyPi: `$ pip install pythran`
- Custom Debian repo: `$ apt-get install pythran`

# We need more peons

- Pythonic needs a serious cleanup
- Typing module needs better error reporting
- OSX support is partial and Windows support is on-going
- `numpy.random` and `numpy.linalg`

# THE END

## **AUTHORS**

Serge Guelton, Pierrick Brunet, Mehdi Amini, Adrien Merlini, Alan Raynaud, Eliott Coyac...

## **INDUSTRIAL SUPPORT**

Silkan, NumScale, →You←

## **CONTRIBUTE**

#pythran on FreeNode, pythran@freelists.org, GitHub repo