




# Cleaning Confused Collections of Characters @ PyDataParis 2015

Ian Ozsvald @IanOzsvald ModellInsight.io

# Who am I?

---

- Past speaker+teacher at PyDatas, EuroPythons, PyCons, PyConUKs
- Co-org of PyDataLondon
- O'Reilly Author
- ModelInsight.io for NLP+ML  
IP creation in London
- “I clean data” #sigh



# Unstructured data->Value

---

- Increasing rate of growth and meant for human consumption
- Hard to:
  - Extract
  - Parse
  - Make machine-readable
- It is also very valuable...part of my consultancy - we're currently automating recruitment:
- Previously - eCommerce recomm., house price & logistics prediction
- Uses: search, visualisation, new ML features
- Most industrial problems messy, not “hard”, but time consuming!
- How can we make it easier for ourselves?

***Elevate***

# What can we extract?

---

- Plain text (languages? platforms? broken files?)
- HTML and XML (e.g. ePub)
- PDFs
- PDF tables
- Image containing text
- (Audio files with speech)

# Extracting from HTML/XML

---

- Assuming you've scraped (e.g. scrapy)
- regular expressions (brittle)
- BeautifulSoup
- xpath via scrapy or lxml
  - `s.xpath('..//span[@class="at_sl"]/text()')[0].extract()`
- *You need unit tests*



# Extracting text from PDFs & Word

---

- pdftotext (Linux), pdfminer (Python 2 with 3 port, 'slow')
- Apache Tika (Java - via jnius?)
- <http://textract.readthedocs.org/en/latest/> (Python)
- Difficulties
  - Formatting probably horrible
  - No semantic interpretation (e.g. CVs)
  - Keyword stuffing, images, out-of-order or multi-column text, tables  
#sigh
- “Content ExtRactor and MINEr” (CERMINE) for academic papers
- Commercial CV parsers (e.g. Sovren)

# Extracting tables from PDFs


---

- ScraperWiki's <https://pdftables.com/> (builds on pdfminer)
- <http://tabula.technology/> (Ruby/Java OS, seems to require user intervention)
- messytables (Python, lesser known, autoguesses dtypes)

# Extracting text from Images

---

- OCR e.g. tesseract (below)
- Abbyy's online commercial service





# Fixing badly encoding text

---

- <http://ftfy.readthedocs.org/en/latest/>
  - The word schön might appear as schÃ¶n.
  - An em dash (–) might appear as â€”.
- HTML unescaping:  
(also ftfy)

```
In [3]: html.unescape('&pound;682m')  
Out[3]: '£682m'
```
- chromium-compact-language-detector will guess human language from 80+ options (so you can choose your own decoding options)

# Interpreting dtypes

---

- Use pandas to get text table (e.g. from JSON/CSV)
- Dates are problematic unless you know their format (next slides), Labix dateutil helpful
- Categories (e.g. “male”/”female”) are easily spotted by eye
- [“33cm”, “22inches”, ...] could be easily converted
- *Could you write a module to suggest possible conversions on dataframe for the user (and notify if ambiguities are present e.g. 1/1 to 12/12...MM/DD or DD/MM)?*

# Date examples

---

```
11 01 2014,10
12 01 2014,11
13 01 2014,12
14 01 2014,13
15 01 2014,14
16 01 2014,15
17 01 2014,16
```


```
In [10]: df_euro.ix[9:16]
Out[10]:
```

	value
date	
2014-01-11	9
2014-01-12	10
2014-01-13	11
2014-01-14	12
2014-01-15	13
2014-01-16	14
2014-01-17	15


- The default is for US-style (MMDD), not Euro-style (DDMM)
- `pd.from_csv(parse_dates=[co/s], dayfirst=False)`

# MMDD (default) vs DDMM

Timeseries with DDMM (dayfirst=True)  
interpretation of Euro data, from 2nd Jan



Timeseries with MMDD (dayfirst=False pandas default)  
interpretation of Euro data



# Merging two data sources

---

- `pd.merge(df1, df2)` # exact keys, SQL-like
- `fuzzywuzzy/metaphone` for approximate string matching
- DataMade's `dedupe.readthedocs.org` to identify duplicates (or OpenRefine)

first name	last name	address	phone
-----			
bob	roberts	1600 pennsylvania ave.	555-0123
Robert	Roberts	1600 Pensylvannia Avenue	



# Manual Normalisation

- Eyeball the problem, solve by hand
- *Lots of unit-tests!*
- `lower()` # “Accenture”->”accenture”
- `strip()` # “ this and ”->”this and”
- `replace(<pattern>,””)` # “BigCo Ltd”->”BigCo”
- `unicode` # “áéîöü”->”aeiou”
- normalise unicode (e.g. >50 dash variants!)
- NLTK stemming & WordNet ISA relt.

Character	Name	Browser	Image
U+002D	HYPHEN-MINUS	-	<a href="#">view</a>
U+058A	ARMENIAN HYPHEN	-	<a href="#">view</a>
U+05BE	HEBREW PUNCTUATION MAQAF	-	<a href="#">view</a>
U+1400	CANADIAN SYLLABICS HYPHEN	=	<a href="#">view</a>
U+1806	MONGOLIAN TODO SOFT HYPHEN	+	<a href="#">view</a>
U+2010	HYPHEN	-	<a href="#">view</a>
U+2011	NON-BREAKING HYPHEN	-	<a href="#">view</a>
U+2012	FIGURE DASH	—	<a href="#">view</a>
U+2013	EN DASH	—	<a href="#">view</a>
U+2014	EM DASH	—	<a href="#">view</a>
U+2015	HORIZONTAL BAR	—	<a href="#">view</a>
U+2E17	DOUBLE OBLIQUE HYPHEN	+	<a href="#">view</a>
U+2E1A	HYPHEN WITH DIAERESIS	˚	<a href="#">view</a>
U+2E3A	TWO-EM DASH	–	<a href="#">view</a>
U+2E3B	THREE-EM DASH	—	<a href="#">view</a>
U+2E40	DOUBLE HYPHEN	⌘	<a href="#">view</a>
U+301C	WAVE DASH	〰	<a href="#">view</a>
U+3030	WAVY DASH	〰	<a href="#">view</a>
U+3040	KATAKANA-HIRAGANA DOUBLE HYPHEN	=	<a href="#">view</a>
U+FE31	PRESENTATION FORM FOR VERTICAL EM DASH		<a href="#">view</a>
U+FE32	PRESENTATION FORM FOR VERTICAL EN DASH		<a href="#">view</a>
U+FE38	SMALL EM DASH	-	<a href="#">view</a>
U+FE03	SMALL HYPHEN-MINUS	-	<a href="#">view</a>
U+FF0D	FULLWIDTH HYPHEN-MINUS	-	<a href="#">view</a>

# Rule lists

---

- Don't forget the old and simple approaches
- Big lists of exact-match rules
- Easy to put into a SQL DB for quick matches!
- A set/dict of strings is super-quick in Python (or use e.g. MarissaTrie)

# Automated Normalisation?

- My annotate.io
- Why not make the machine do this for us?

What you have:

"To 53K w/benefits"

"30000 OTE plus bonus"

"£55000 salary"

"Forty two thousand GBP"

What you want:

"53000"

"30000"

"55000"

"42000"



# Machine Learn the rules?


---

- What if we extend dedupe's idea?
- Can we give examples of e.g. company names that are similar and generalise a set of rules for unseen data?
- Could we train given a small set of data and re-train when errors occur on previously unseen data?

# Visualising new data sources

---

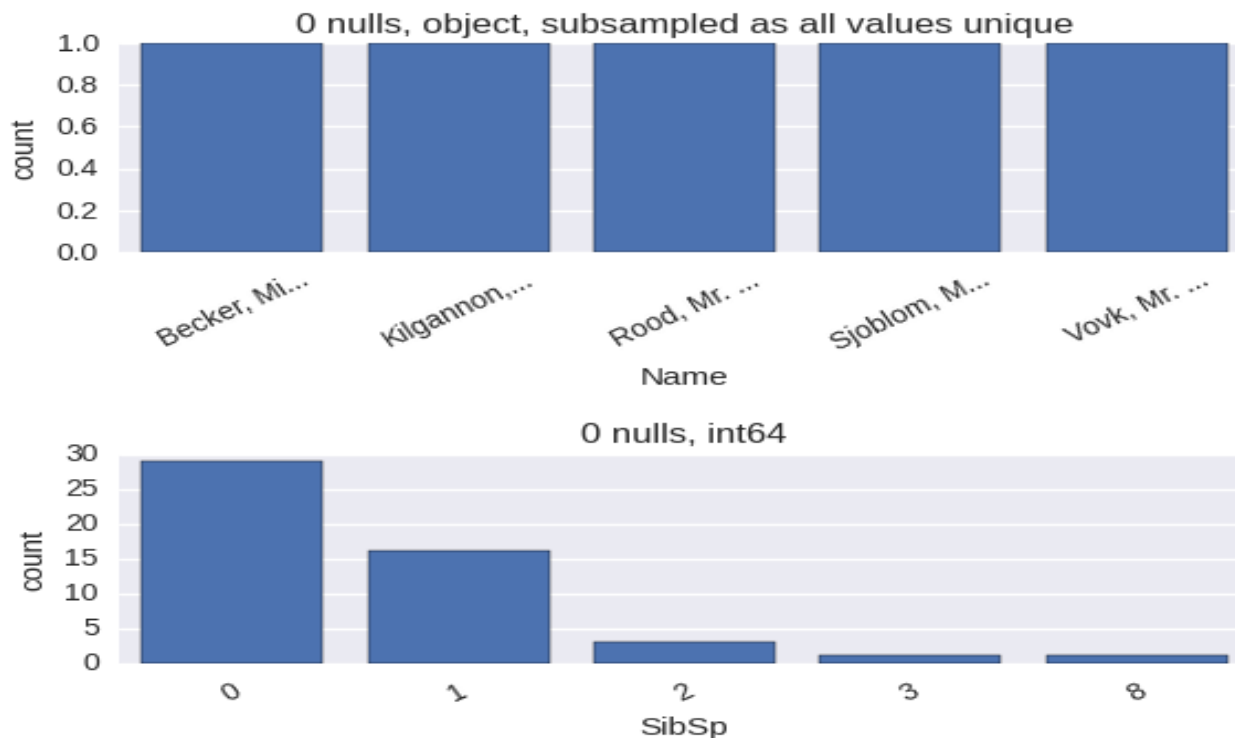
- [setosa.io/csv-fingerprint/](http://setosa.io/csv-fingerprint/)
- SeaBorn (or Bokeh?)
- Do you have good tools?





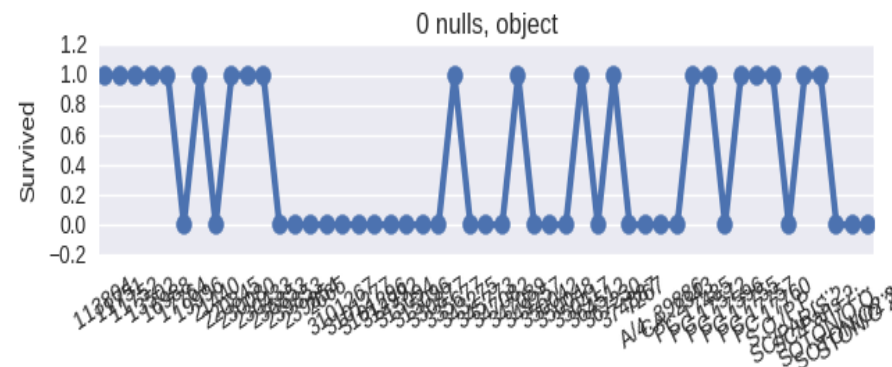
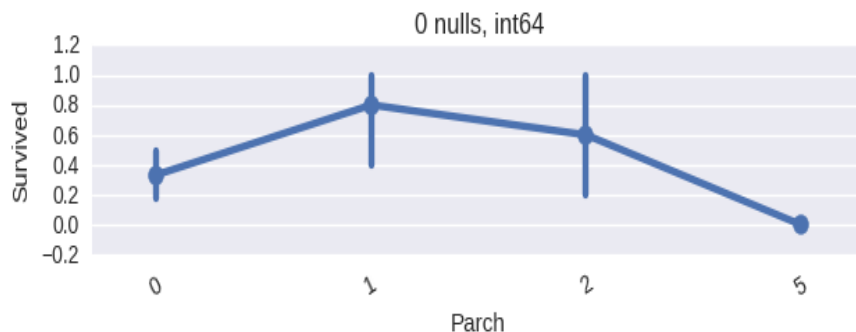

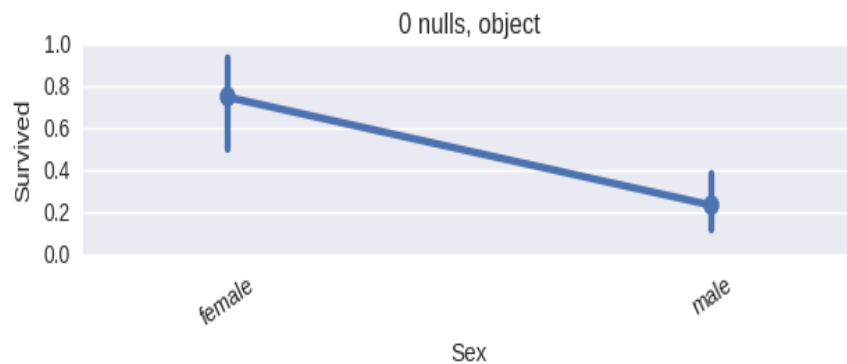
# Visualising new data sources

- [github /ianozsvald/dataframe\\_visualiser](#)



# Visualising new data sources

- [github /ianozsvald/dataframe\\_visualiser](https://github.com/ianozsvald/dataframe_visualiser)



# Automate feature extraction?

---

- Can we extract features from e.g. Id columns (products/rooms/categories)?
- We could identify categorical labels and suggest Boolean column equivalents
- We could remove some of the leg-work...you avoid missing possibilities, junior data scientists get “free help”
- What tools do you know of and use?

# Getting started

---

- Cleaning is more R&D than engineering
- FACT: Your data has missing items + it lies
- Visualise it
- Set realistic milestones, break into steps, have lots of tests
- Have a gold standard for measuring progress
- Aim for a high-quality output

# Tips

---

- Lots of lesser-known good tools out there!
- APIs like Alchemy + DBPedia for Entity Recognition and Sentiment Analysis
- Python 3.4+ makes Unicode easier
- USE: pandas, StackOverflow



# Closing...

---

- Give me feedback on [annotate.io](http://annotate.io)
- *Give me your dirty-data horror stories*
- <http://ianozsvald.com/>
- PyDataLondon monthly meetup
- PyDataLondon conference late June(?)
- *Do you have data science deployment stories for my keynote at PyConSweden? What's "hardest" in data science for your team?*